Remote GPU inference using SOL
ROS2 (Fast DDS TCP / CycloneDDS UDP) vs vAccel, 1000 inputs per run

Latency (top row)
- Network + Framework Overhead
- Pre/Post-processing
- Inference

ROS2 Fast DDS (TCP) | ROS2 CycloneDDS (UDP) | vAccel

Energy (bottom row)
- Edge GPU
- Edge CPU
- Robot CPU

Image classification

Median latency per input (ms)

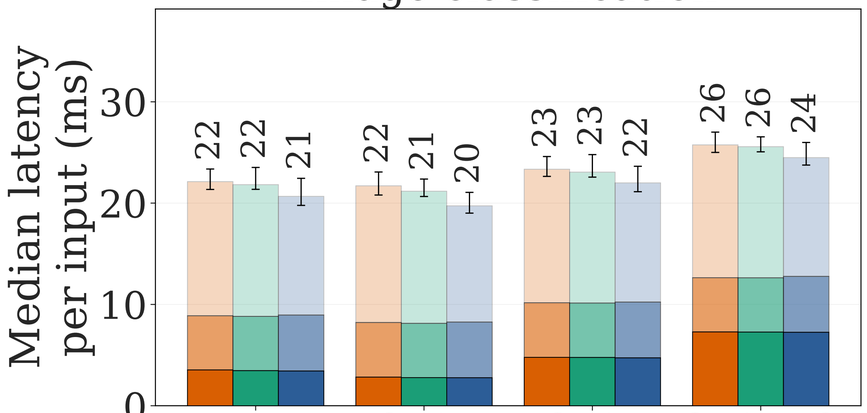

Total energy (kJ)

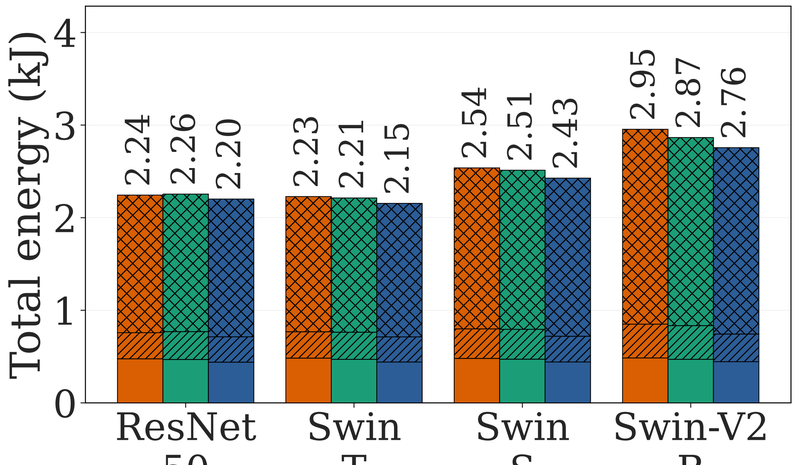

Video classification

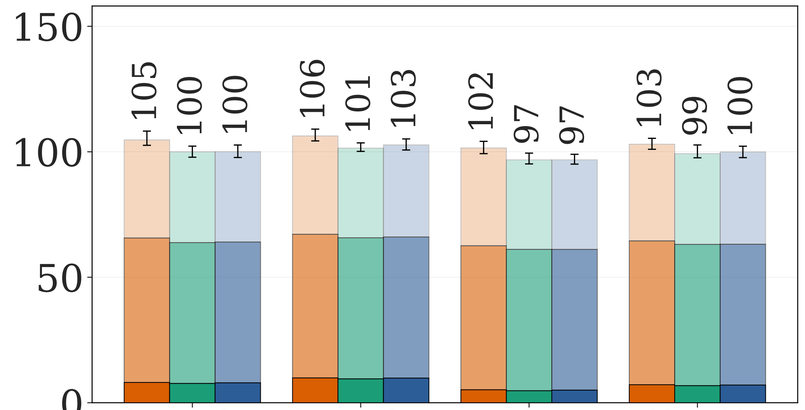

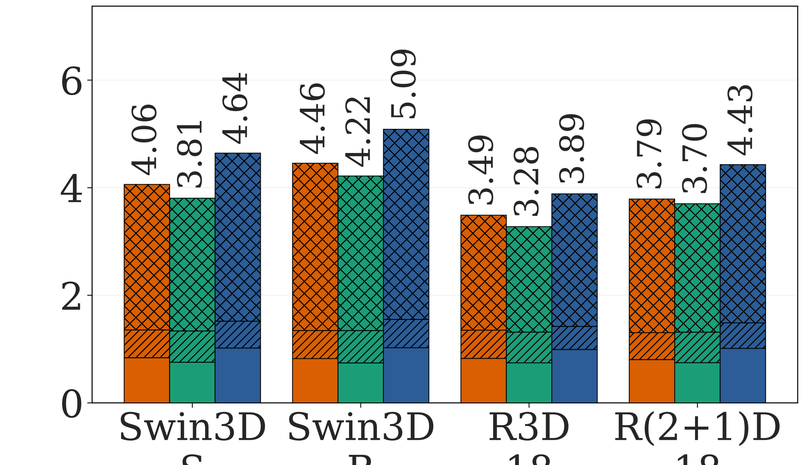
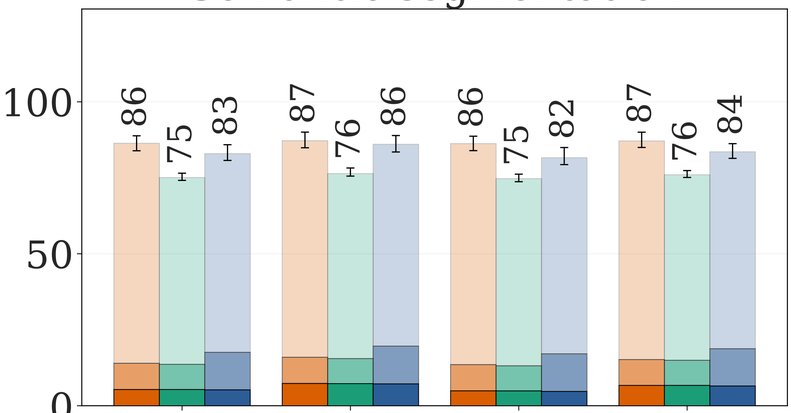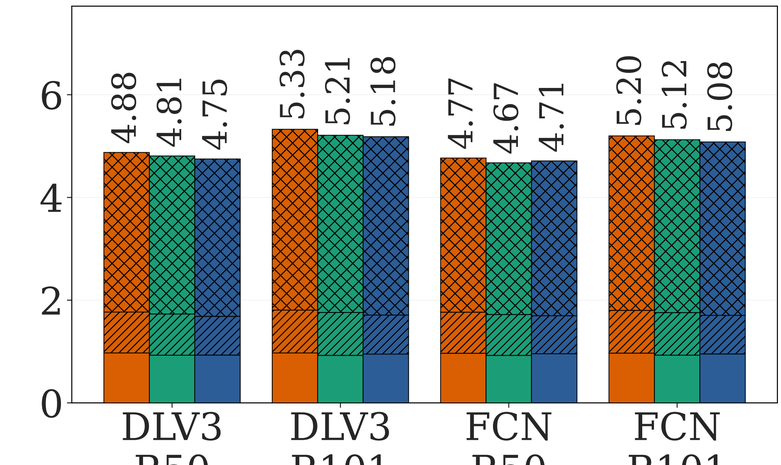
Semantic segmentation