

Robot-side inference latency under local execution and edge offloading

